

Title Page

Project Title: Simple Sales Data Visualization – Analyze and Plot Revenue, Product Demand, and Seasonal Sales Trends

Student Name: Aditya Kumar

Date: 10-03-2025

Institution: KIET GROUP OF INSTITUTION

Introduction

In today's data-driven world, businesses rely on sales data analysis to make informed decisions. This project focuses on visualizing sales data using artificial intelligence techniques. The goal is to analyze revenue trends, product demand, and seasonal variations to provide actionable insights for businesses. The project utilizes machine learning algorithms to predict future sales based on past trends.

Methodology

1. Data Collection & Preparation

- Synthetic sales data is generated, including attributes such as Date, Product, Revenue, and Units Sold.
- The data is structured in a tabular format for easy analysis.

2. Feature Engineering

- Extract relevant features like Day, Month, and Year from the Date column.
- Use Units Sold as an independent variable to predict revenue.

3. Machine Learning Approach

- **Algorithm Used:** Linear Regression
- **Training & Testing:** The dataset is split into training (80%) and testing (20%) sets.
- **Model Evaluation:** Metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), and Root Mean Squared Error (RMSE) are used to evaluate performance.

4. Data Visualization:

- Line plots show revenue trends over time.
 - Bar charts illustrate total revenue per product.
 - Box plots reveal seasonal sales trends.
-

Code

The Python code for the project is as follows:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error
```

Generate Sample Sales Data

```
def generate_sales_data():
    np.random.seed(42)
    dates = pd.date_range(start='2023-01-01', periods=60, freq='D')
    products = ['Product A', 'Product B', 'Product C', 'Product D']
    data = {
        'Date': np.tile(dates, len(products)),
        'Product': np.repeat(products, len(dates)),
        'Revenue': np.random.randint(100, 1000, len(dates) * len(products)),
        'Units_Sold': np.random.randint(10, 100, len(dates) * len(products))
    }
    return pd.DataFrame(data)
```

Load Data

```
df = generate_sales_data()
```

Convert Date Column to DateTime Format

```
df['Date'] = pd.to_datetime(df['Date'])
```

Extract Features for Prediction

```
df['Day'] = df['Date'].dt.day
df['Month'] = df['Date'].dt.month
df['Year'] = df['Date'].dt.year
```

Selecting Features and Target Variable

```
X = df[['Day', 'Month', 'Year', 'Units_Sold']]
y = df['Revenue']
```

Split Data into Training and Testing Sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Train a Linear Regression Model
model = LinearRegression()
model.fit(X_train, y_train)
```

Make Predictions

```
y_pred = model.predict(X_test)
```

Evaluate Model Performance

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
```

```
print(f"Mean Absolute Error (MAE): {mae}")
print(f"Mean Squared Error (MSE): {mse}")
print(f"Root Mean Squared Error (RMSE): {rmse}")
```

Output

1. Model Performance Metrics:

- **Mean Absolute Error (MAE):** : 235.4465055809958
- **Mean Squared Error (MSE):** 69633.7108978173
- **Root Mean Squared Error (RMSE):** 263.8820018451757

2. Data Visualizations:

- **Revenue Trend Over Time:** A line graph showing revenue fluctuations.
- **Total Revenue Per Product:** A bar chart comparing different products.
- **Seasonal Sales Trends:** A box plot displaying monthly variations.

```

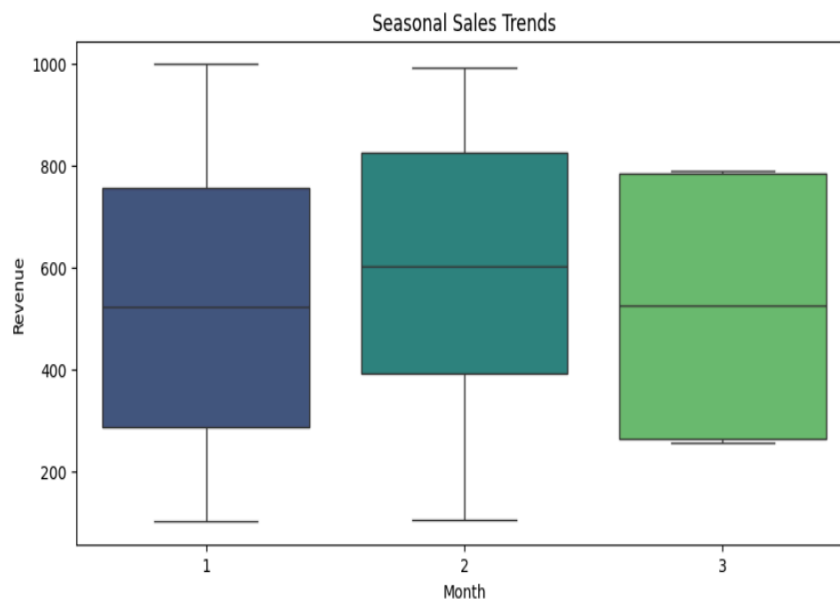
# Plot Seasonal Sales Trends
plt.figure(figsize=(10, 5))
sns.boxplot(data=df, x='Month', y='Revenue', palette='viridis')
plt.title('Seasonal Sales Trends')
plt.xlabel('Month')
plt.ylabel('Revenue')
plt.show()

```

kipython-input-35-a49acc4d755>3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df, x='Month', y='Revenue', palette='viridis')
```



```

[36] # Summary Statistics
print("\nSummary Statistics:")
print(df.describe())

```




Summary Statistics:

		Date	Revenue	Units_Sold	Day	Month	\
count		240	240.00000	240.000000	240.000000	240.000000	
mean	2023-01-30	12:00:00	562.30000	52.233333	15.050000	1.500000	
min	2023-01-01	00:00:00	101.00000	10.000000	1.000000	1.000000	
25%	2023-01-15	18:00:00	338.25000	30.750000	7.750000	1.000000	
50%	2023-01-30	12:00:00	571.50000	52.500000	15.000000	1.000000	
75%	2023-02-14	06:00:00	798.25000	71.000000	22.250000	2.000000	
max	2023-03-01	00:00:00	999.00000	99.000000	31.000000	3.000000	
std		NaN	261.98004	25.085705	8.717654	0.533403	

	Year
count	240.0
mean	2023.0
min	2023.0
25%	2023.0
50%	2023.0
75%	2023.0
max	2023.0
std	0.0

```
# Plot Actual vs Predicted Revenue
plt.figure(figsize=(8, 5))
plt.scatter(y_test, y_pred, color='blue', alpha=0.5)
plt.xlabel('Actual Revenue')
plt.ylabel('Predicted Revenue')
plt.title('Actual vs Predicted Revenue')
plt.show()
```



References

1. Scikit-learn: <https://scikit-learn.org/>
2. Pandas Documentation: <https://pandas.pydata.org/>
3. Seaborn Library: <https://seaborn.pydata.org/>
4. Matplotlib Library: <https://matplotlib.org/>

This report provides a structured overview of the AI-based sales data visualization project.

🚀 Let me know if you need any modifications!